

How JVM works

Outline

- Frames - space for calculation
~ operand stack, local variables, ...
- Bytecode instructions
- Method calls create frames
- Heap for objects
- Symbolic references to methods

Frames

or stack frames

Each time you call a method in java, it gets a new frame constructed for it.

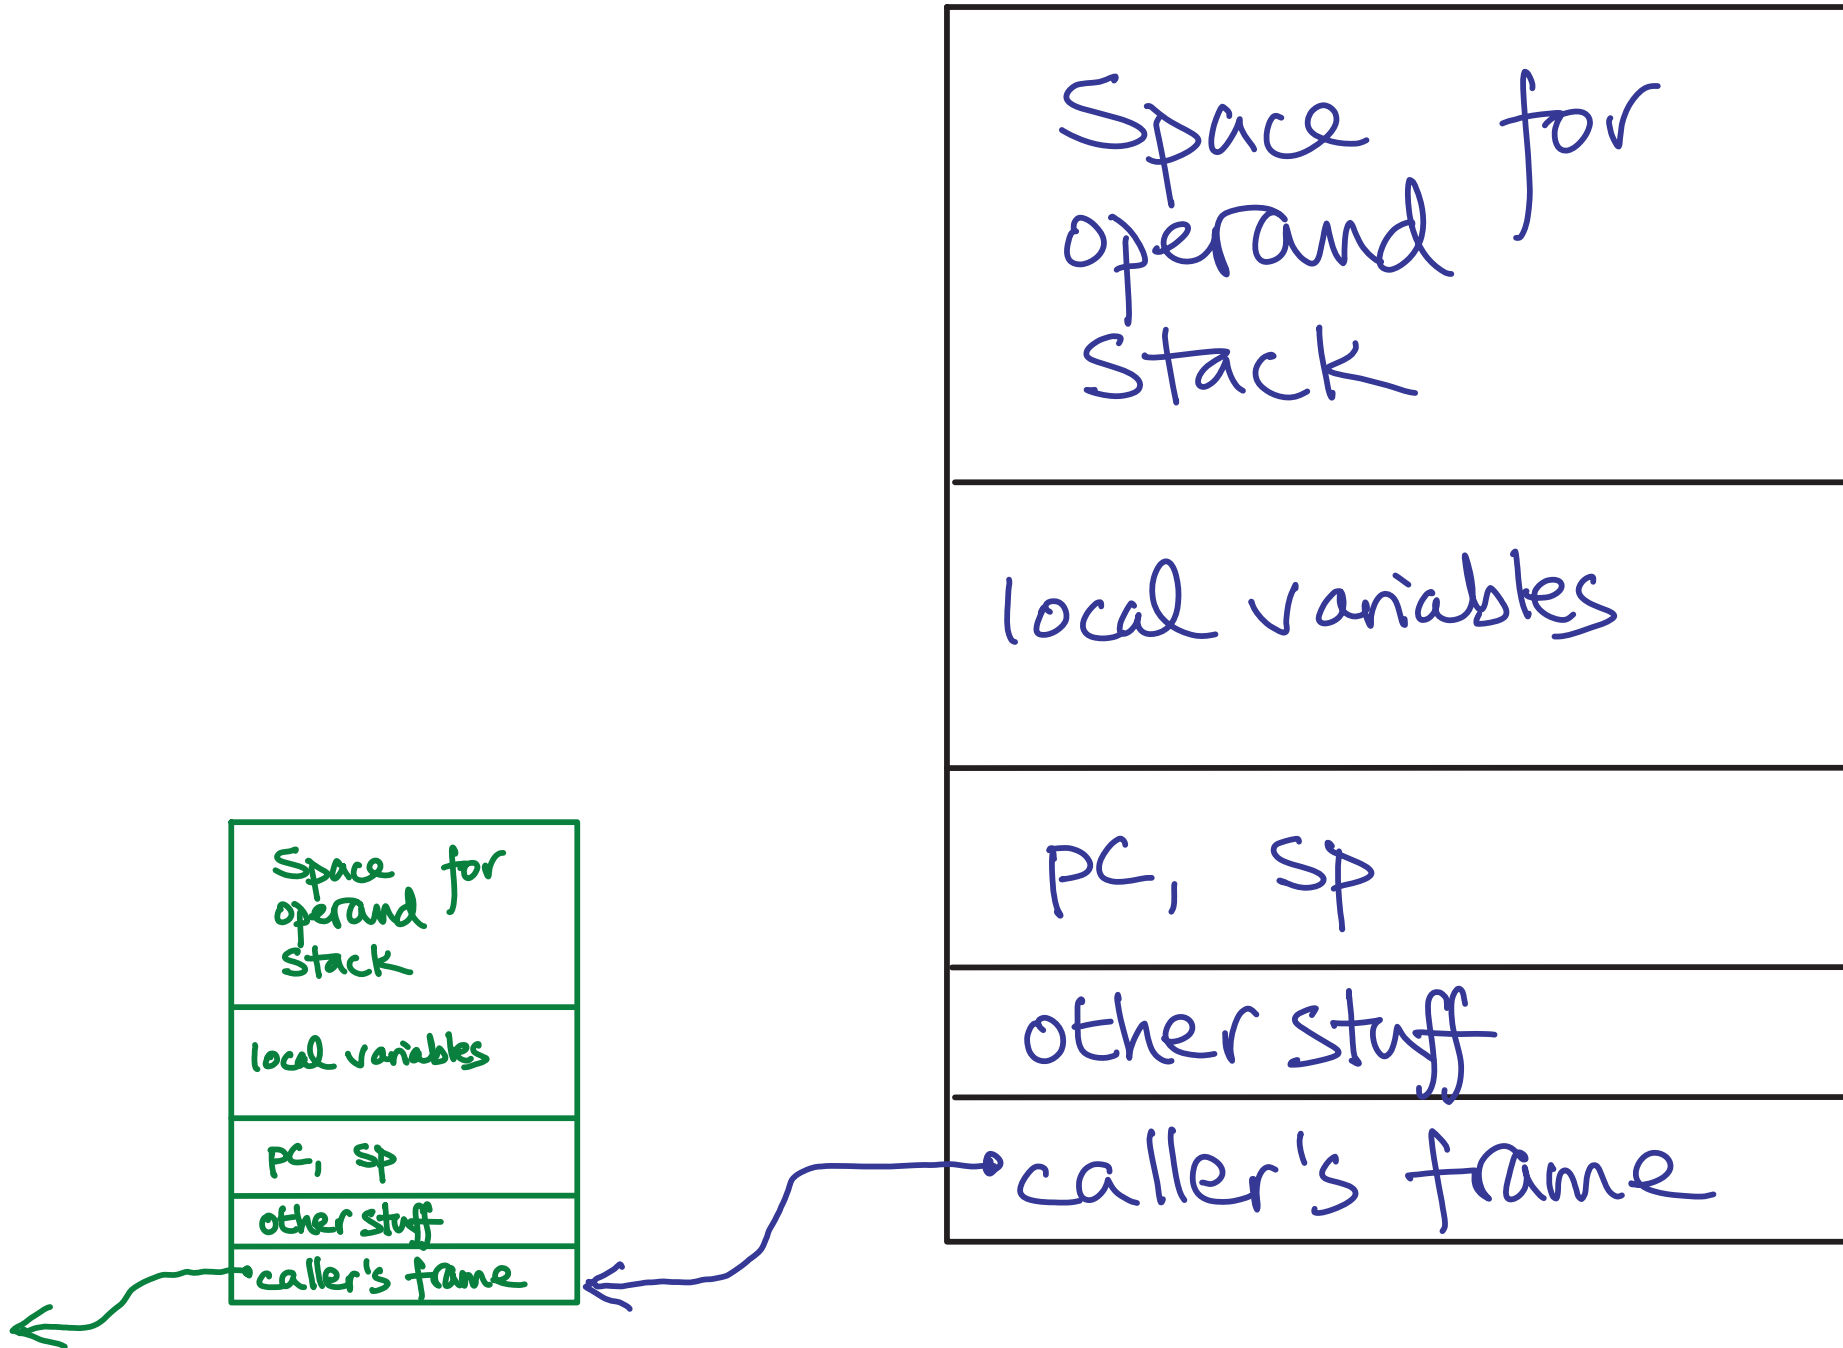
The frame has -

same principle as registers

- storage space for local variables
- space for an operand stack
- program counter & stack pointer
- reference back to frame of caller
- other stuff

for virtual machine, not CPU

Frame



Entries are 4 bytes each

on operand stack or local variables

- enough for bool^{1 bit}, byte^{1 byte}, short², int⁴, char², float⁴
- also enough for any reference⁴ value
- for long⁸, double⁸ need two consecutive entries

What is a local variable?

- ① "this" (for a non-static method)
 - reference to the object "this"
 - Ropt as local variable with index 0
- ② parameters of method
Indexes start 0 (static) or 1 (non-static)
- ③ Variables declared in method

Sometimes "local variable" specifically means ③

What is not a local variable?

- ① Instance variables
(non-static fields)
- ② Class variables
(static fields)

} both
declared
outside
method

How many local variables does
this constructor have?

```
public class PosVal {  
    private static int nextSerial = 0;  
    private int serial;  
    private int val; //invariant: val >= 0  
  
    public PosVal(int initVal){  
        int v = initVal;  
        if (v < 0){  
            v = -v;  
        }  
        val = v;  
        serial = nextSerial;  
        nextSerial += 1;  
    }  
}
```

What is a local variable?

- ① "this" (for a non-static method)
 - reference to the object "this"
 - kept as local variable with index 0
- ② parameters of method
Indexes start 0 (static) or 1 (non-static)
- ③ Variables declared in method

How many local variables does
the constructor have?

THREE

```
public class PosVal {  
    private static int nextSerial = 0;  
    private int serial;  
    private int val; //invariant: val >= 0
```

```
    public PosVal(int initVal) {  
        int v = initVal;  
        if (v < 0) {  
            v = -v;
```

```
        }  
        val = v;  
        serial = nextSerial;  
        nextSerial += 1;  
    }
```

instance
variables

class variable

① "this"

① parameter


②

↑
indexes
= slot numbers

Local variables in JVM

JVM doesn't know the names of the local variables (from Java source).

- Refers to them by slot numbers starting at 0



- For non-static method, variable 0 is this
Then parameters start at slot 1,
then other local variables
- For static method, parameters start at slot 0
- For long, double use number of first of pair of slots
8 bytes = 2 slots

Can operand stack ever overflow
(run out of space)?

No.

- Compiler works out exactly how much space needed
- Loader checks each method to verify no overflow or underflow possible

StackOverflowError is different -
a new frame is needed, but there's
not enough memory.

Bytecode instructions

~ 256 different ones!
I shan't describe them all

Each instruction has at least one byte

— opcode

May have more operand(s)

Two separate meanings of "operand" — here, & entry on operand stack

A single operand may be 2 or more bytes together as an integer. Then "bigendian" — most significant bytes come first.

For each opcode there's a human-readable mnemonic

Arithmetic on operand stack

e.g. add - adds top two stack entries

..., val1, val2 \Rightarrow ..., val1 + val2

BUT different opcodes for different types

mnemonic	opcode (hex)	type	
iadd	60	int	i
ladd	61	long	l
fadd	62	float	f
dadd	63	double	d

Similarly for other operations.

Also sometimes b - byte s - short
 a - address (reference)

Errors?

What happens if ---

- use fadd on int entries?
- there's only one entry on stack?
- other obvious mistakes?

stack
underflow

No checks when operation is executed

BUT JRE verifies code when it loads a class

- checks types used consistently etc

Safeguard against security holes

Pushing constants on the stack

push N ← one operand byte
... → ..., N

b for byte - 1 byte operand

↓
b i push one operand byte
↑

i for int - pushes 4 bytes

"sign extends" operand to 4 bytes
pushes result on stack

Similarly sipush two operand bytes
↑
short

Simple opcodes for common constants

e.g. for int

7 different opcodes {
 iCONST_m1
 iCONST_0
 iCONST_1
 ⋮
 iCONST_5
} no operand needed

... ⇒ ... ,
 -1
 0
 ⋮
 5

pushes
-1, 0, 1, 2, 3, 4 or 5

Load

.. push variable

load variable onto stack

e.g. iload slot number

1-byte operand

push int local variable at given slot
onto stack

Similarly lload, fload, dload, aload
long float double ref (address)

JRE verifier: must use types consistently
e.g. can't load as integer then use as address

It's Java, not C++!

1-byte loads

Special opcodes with no operand for
slot numbers 0, 1, 2, 3

e.g. iload_0, iload_1, iload_2, iload_3

Store

Reverse of load: pops top of operand stack
into variable

e.g. istore slot number
astore_2

Jumps

Must be within current method.
Can't jump to a different method.

Unconditional

goto 2 byte offset

operand is added to address of
goto opcode to give address of next
opcode to execute

operand is
an offset

Similarly: goto-w 4 byte offset
(wide)

Conditional jumps

As before,
but using offsets
for operand N.

Also, e.g.

if-icmp [↑]
int comparison

Conditional jumps - e.g.
ifeq N - jumps to N if val = 0

..., val → ...
if-~~compare~~cmpeq N - jumps to N if val1 = val2
..., val1, val2 → ...

jump if val $\left\{ \begin{array}{l} = \\ < \\ \leq \\ \neq \\ > \\ \geq \end{array} \right\} 0$ eq
lt
le
ne
gt
ge

6 operators ifeq, iflt, ifle etc;
also if-cmpeq, if-cmplt, etc.

Call and return

two parameters

Say method B calls $A(p_0, p_1)$

& A returns a result

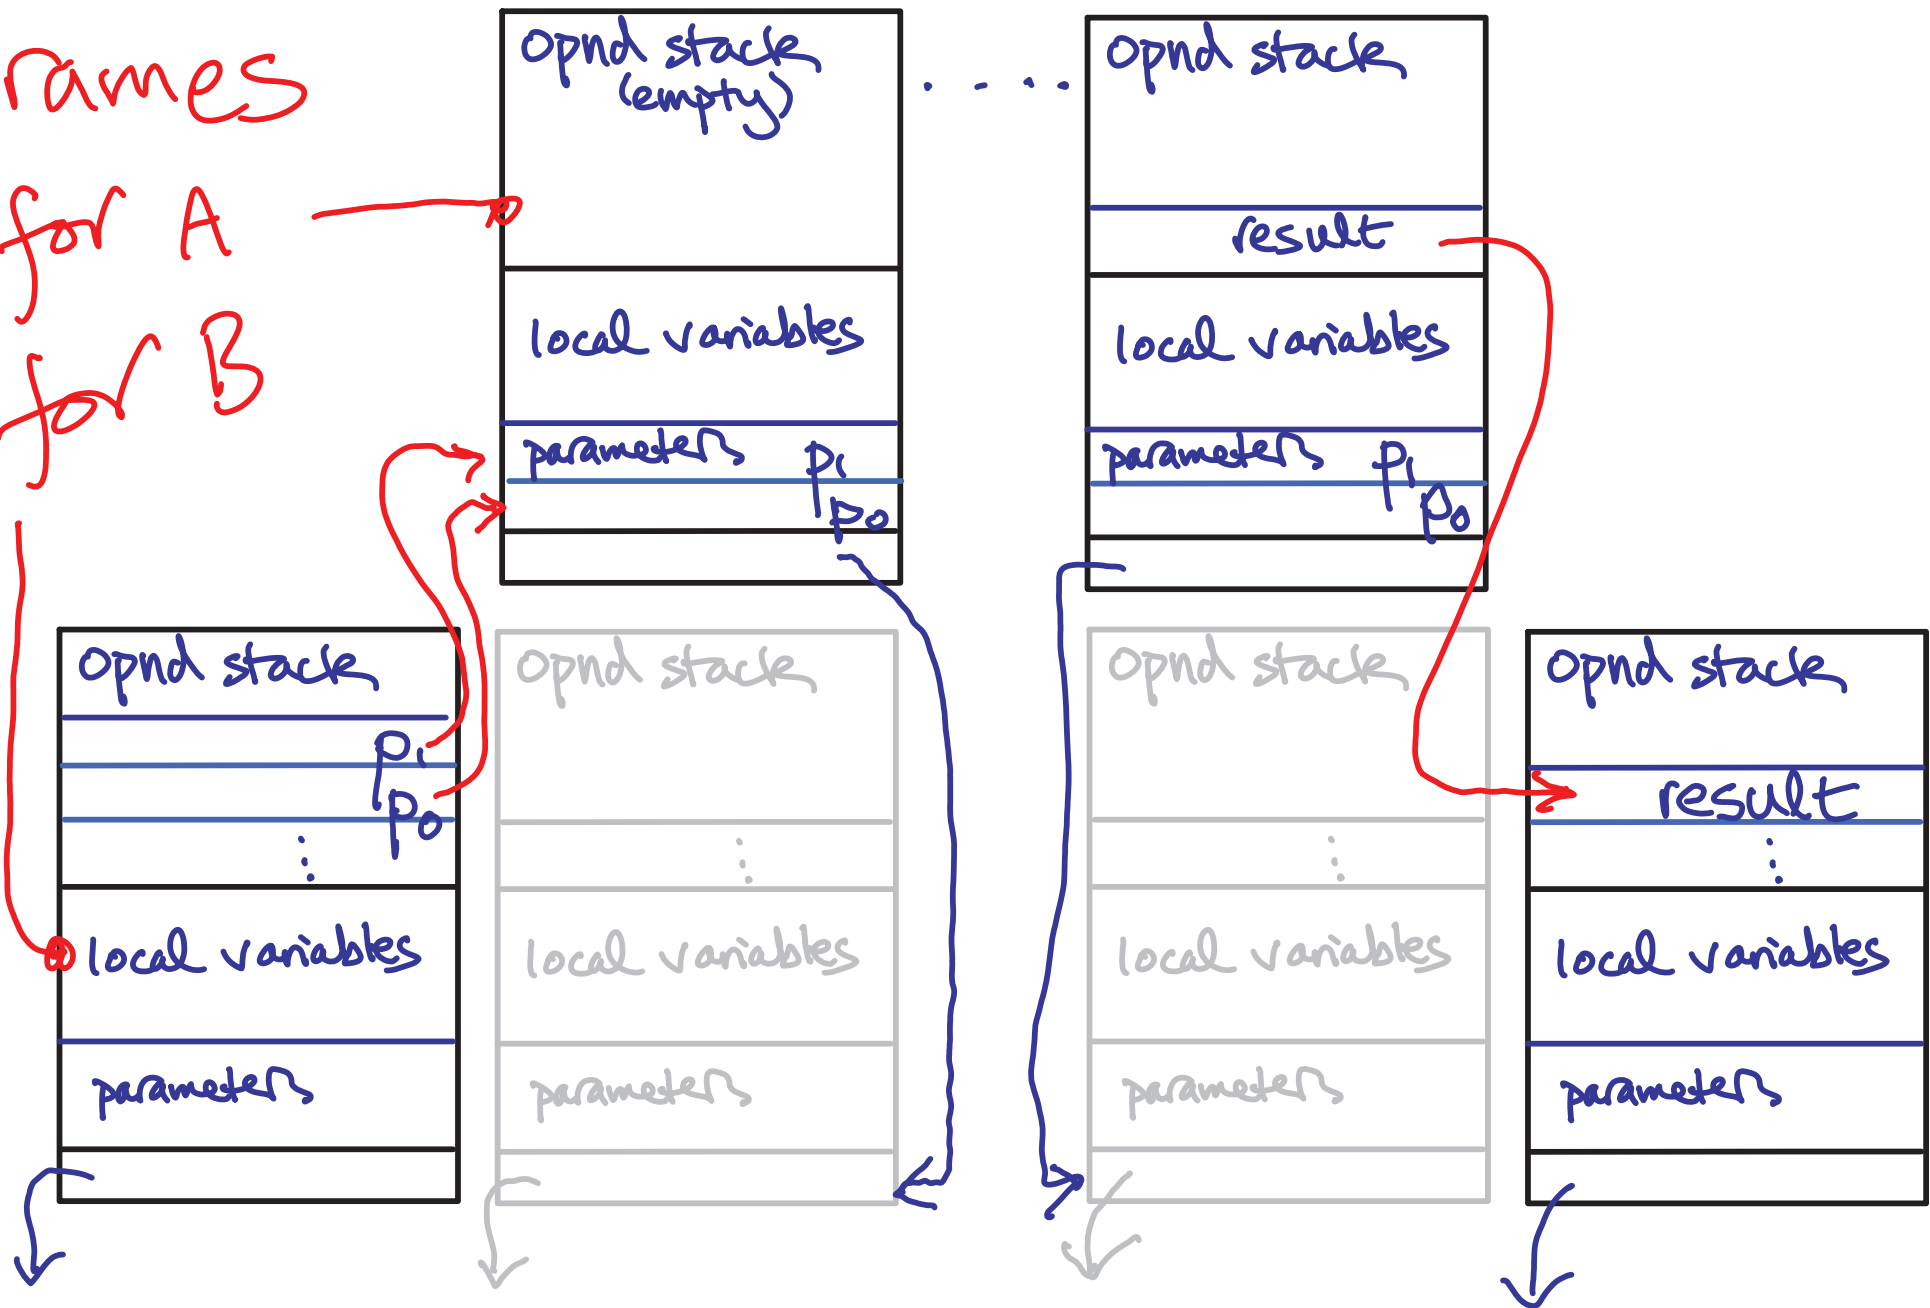
not void

- ① B calculates actual parameters on its operand stack
 - ② Construct stack frame for A, with parameters taken from B's opnd stack
 - ③ A calculates its result on its opnd stack
 - ④ result transferred to B's opnd stack
 - ⑤ Return to B, throw away A's stack frame
- On B's operand stack: $A(p_0, p_1)$ has effect of
- $\dots, p_0, p_1 \rightarrow \dots, \text{result}$

B calls A ... A executes ... A returns to B

Frames

for A
for B



Current frame — for method being executed

- While A is executed, its frame is current
 - B's frame is not. Greyed out on previous slide

- When A returns, its frame is destroyed
 - B's frame becomes current again

Saving return addresses

- Each frame has its own pc
- While A is executed, its pc is used
- When A returns, B resumes with its old pc value
- B's local variables are also unchanged by A.

Linked frames have the effect of a return stack.

In fact, the chain of linked frames is officially called a stack in JVM